

PROJECT REPORT

Explainable Multiclass Sleep Disorder Classification

An End-to-End Machine Learning Pipeline for Predicting Insomnia and Sleep Apnea from Lifestyle and Physiological Health Data

Supervised Learning • Class-Imbalance Handling • Explainable AI (XAI)

Prepared by Rushikesh Khamgaonkar | August 2026

Final Model	Test Accuracy	Test Macro F1	Dataset Size
XGBoost (Tuned)	93.50%	88.50%	10,000 records

Dataset: Sleep Health & Lifestyle Records (10,000 individuals, 13 features) • 3 Target Classes: No Sleep Disorder • Insomnia • Sleep Apnea
Leakage-Free Pipeline • Stratified Cross-Validation • SHAP-Based Model Interpretability • Flask Deployment

Table of Contents

1. Executive Summary	3
1.1 What This Project Demonstrates	3
2. Project Overview	4
2.1 Project Objective	4
2.2 Business Objective	4
2.3 Problem Statement	4
2.4 Dataset Description	4
3. Data Preparation & Understanding	5
3.1 Initial Data Preparation	5
3.2 Data Quality & Validation Checks	5
3.3 Descriptive & Target Distribution Summary	5
4. Exploratory Data Analysis (EDA)	6
4.1 Target Variable Distribution	6
4.2 Numerical Feature Distribution	6
4.3 Categorical Feature Distribution	7
4.4 Feature Relationships with the Target	8
4.5 Correlation Analysis	9
4.6 Multivariate Relationship Analysis	10
4.7 Outlier Analysis	10
5. Feature Engineering	12
5.1 Engineered Features	12
5.2 Engineered Feature Validation	12
6. Data Preprocessing & Train-Test Split	12
6.1 Preprocessing Pipeline	12
6.2 Stratified Train-Test Split	12
7. Model Development & Selection	13
7.1 Modeling Approach	13
7.2 Baseline Cross-Validation Results	13
7.3 Class Weight vs. SMOTE Comparison	14
7.4 Hyperparameter Tuning	14
8. Final Model Evaluation	15
8.1 Held-Out Test Set Performance	15
8.2 Confusion Matrix	16
8.3 Multiclass ROC-AUC	16
8.4 Precision-Recall Curves & PR-AUC	17
8.5 Class-wise Performance	17
9. Model Explainability (XAI)	18
9.1 SHAP Global Feature Importance	18
9.2 SHAP Feature Interaction Analysis	19
9.3 Explainability Summary	20
10. Conclusion & Key Takeaways	20
10.1 Summary of Findings	20
10.2 Final Recommendation	21
10.3 Business Impact & Recommendations	21
10.4 Skills & Techniques Demonstrated	21

1. Executive Summary

This report documents an end-to-end machine learning project that classifies an individual's sleep health status using demographic, lifestyle, and physiological data. The system distinguishes between three outcomes — No Sleep Disorder, Insomnia, and Sleep Apnea — from a dataset of 10,000 patient records spanning occupation, activity, stress, sleep quality, and cardiovascular indicators.

The project follows a leakage-aware, eight-phase pipeline: data preparation, data understanding, exploratory analysis, feature engineering, preprocessing, model development, hyperparameter tuning, and final evaluation with explainability. Five classification algorithms — Logistic Regression, Decision Tree, Random Forest, Support Vector Machine, and XGBoost — were each benchmarked under two class-imbalance strategies (class weighting and SMOTE) using Stratified 5-Fold Cross-Validation, since the target is imbalanced at roughly 78% / 12% / 10% across its three classes.

Headline Result

A tuned XGBoost classifier was the best-performing model, achieving 93.50% accuracy and a Macro F1-score of 88.50% on a completely unseen, held-out test set — outperforming Logistic Regression, Decision Tree, Random Forest, and SVM under both imbalance-handling strategies tested.

Beyond raw predictive accuracy, the project prioritised clinical interpretability. SHAP (SHapley Additive exPlanations) was applied to the final XGBoost model to quantify which features drive each prediction, confirming that Heart Rate, Sleep Duration, and the engineered Sleep Quality Index are the dominant physiological and behavioural signals, while demographic and occupational attributes contribute comparatively little.

1.1 What This Project Demonstrates

- End-to-end ML pipeline design — from raw Excel ingestion to a deployable, explainable classifier.
- Imbalanced-classification discipline — stratified splitting, class weighting, and SMOTE, compared side-by-side rather than assumed.
- Leakage-safe engineering — preprocessing and resampling refit independently inside every cross-validation fold.
- Feature engineering for health data — pulse pressure, a composite sleep-quality index, and a stress-sleep interaction term, each derived without touching the target.
- Breadth across model families — linear, tree-based, kernel-based, and gradient-boosted classifiers.
- Explainable AI (XAI) — global and interaction-level SHAP analysis, turning a 'black-box' gradient-boosted model into a transparent, stakeholder-readable decision system.

2. Project Overview

2.1 Project Objective

Develop an explainable multiclass classification system that predicts an individual's sleep-disorder status — No Sleep Disorder, Insomnia, or Sleep Apnea — from demographic, occupational, lifestyle, and physiological attributes, while benchmarking multiple algorithms and imbalance-handling strategies to identify the most reliable, generalisable model.

2.2 Business Objective

- Flag individuals at risk of a sleep disorder from routinely available health and lifestyle data.
- Benchmark class-imbalance strategies and select the strongest model on class-aware metrics, not raw accuracy alone.
- Explain predictions with SHAP to support clinical trust and stakeholder transparency.
- Package the final model as a deployable Flask web application for real-time inference.

2.3 Problem Statement

Sleep disorders such as Insomnia and Sleep Apnea are frequently under-diagnosed despite measurable ties to lifestyle and physiological signals — stress, activity level, blood pressure, and heart rate. Given an individual's demographic, lifestyle, and physiological attributes, this project predicts sleep-disorder status as one of three classes, using a dataset that is naturally imbalanced (78% / 12% / 10%), which makes naïve accuracy an unreliable evaluation metric and motivates the model-selection strategy used throughout this report.

2.4 Dataset Description

The dataset consists of synthetically generated patient-level sleep, lifestyle, and cardiovascular records.

Attribute	Value
Raw Records	10,050 rows (50 exact duplicates removed)
Modelling Records	10,000 records × 13 features
Target Variable	Sleep Disorder (No Sleep Disorder · Insomnia · Sleep Apnea)
Feature Groups	Demographic, Occupational, Lifestyle, Physiological, Cardiovascular
Missing Values	12 of 13 features affected (up to ~5.4% per feature)
Duplicate Records	50 identified and removed
Source	Sleep Health and Lifestyle Dataset (Kaggle, synthetically generated)

Data cleaning note: the combined Blood Pressure field (e.g. "126/83") was split into separate Systolic BP and Diastolic BP columns to enable numerical analysis. The literal string "None" in the target column was explicitly preserved as a valid category during ingestion (rather than parsed as a missing value) and relabelled to "No Sleep Disorder", and inconsistent categorical entries (e.g. "FEMALE", "female", "Female ") were standardised prior to encoding.

3. Data Preparation & Understanding

Before analysis, the raw dataset was cleaned and validated to ensure every downstream statistic and model is computed on a trustworthy, consistent table.

3.1 Initial Data Preparation

- Identifier removal: non-predictive identifier columns (e.g. Person ID) were dropped before analysis.
- Duplicate handling: 50 exact duplicate records were identified and removed, leaving 10,000 unique rows.
- Categorical standardisation: inconsistent casing and stray whitespace in Gender, Occupation, and BMI Category were normalised to a single canonical spelling per category.
- Missing-value standardisation: blank strings were explicitly converted to NaN so that numerical and categorical imputers could detect them consistently.
- Blood Pressure decomposition: the combined "Systolic/Diastolic" text field was split into two numeric columns, Systolic BP and Diastolic BP.

3.2 Data Quality & Validation Checks

Following the initial preparation steps, the dataset was re-validated to confirm the cleaning stage produced a leakage-free, analysis-ready table.

Check	Result
Duplicate records remaining	0
Features with missing values	12 of 13 (imputed later inside the pipeline, not dropped)
Maximum missing-value rate	~5.4% (single feature)
Target class integrity	"None" correctly preserved and relabelled — not read as null
Categorical consistency	Standardised (no duplicate-meaning categories remain)

3.3 Descriptive & Target Distribution Summary

Descriptive statistics confirmed plausible physiological ranges across the dataset (e.g. Heart Rate roughly 50–125 bpm, Sleep Duration roughly 4.5–9.5 hours), with no impossible or out-of-range values detected. The target variable, however, is meaningfully imbalanced — a fact that shapes the modelling strategy for the remainder of this report.

Class	Count	% of Dataset
No Sleep Disorder	7,813	78.13%
Insomnia	1,237	12.37%
Sleep Apnea	950	9.50%

Key takeaway: a model that always predicts "No Sleep Disorder" would already score ~78% accuracy while never identifying a single at-risk patient. This motivated Macro F1 as the primary model-selection metric, alongside class weighting and SMOTE as imbalance-handling strategies.

4. Exploratory Data Analysis (EDA)

Before modelling, the cleaned dataset was explored visually and statistically to understand the shape of the target, feature distributions, categorical breakdowns, correlations, and outliers.

4.1 Target Variable Distribution

The target variable shows a clear three-way class imbalance: healthy individuals dominate the dataset, while Insomnia and Sleep Apnea together make up roughly one in five records.




Figure 4.1 — Distribution of the Sleep Disorder Target Variable

Observation: the 78% / 12% / 10% split confirms the dataset is imbalanced rather than merely skewed, reinforcing the decision to evaluate models on Macro F1 and Balanced Accuracy rather than accuracy alone.

4.2 Numerical Feature Distribution

Histograms with kernel-density overlays were plotted for all nine numerical features to inspect shape, skew, and multimodality ahead of modelling.




Figure 4.2 — Distribution of Numerical Features

Insights: Age and Sleep Duration are broadly bell-shaped; Stress Level and Quality of Sleep are discrete, multi-modal ordinal scales rather than continuous measures; Daily Steps and Systolic BP show mild right-skew consistent with a small subset of more sedentary, higher-risk individuals.

4.3 Categorical Feature Distribution

Gender, Occupation, and BMI Category were profiled to understand demographic composition and category cardinality ahead of one-hot encoding.




Figure 4.3 — Distribution of Categorical Features

Insights: Gender is close to evenly split (4,929 vs. 4,843). Occupation is dominated by Software Engineers, Teachers, and Nurses, with several low-frequency categories (e.g. Scientist, Student) that require careful handling to avoid sparse one-hot columns. Just under half of individuals fall in the Normal BMI category, with Overweight the next-largest group.

4.4 Feature Relationships with the Target

Each numerical feature was compared across the three target classes using box plots to identify which signals separate healthy individuals from those with a disorder.




Figure 4.4a — Numerical Features vs. Sleep Disorder

Insights: Sleep Apnea patients show markedly higher Systolic BP, Diastolic BP, and Heart Rate than the other two classes, alongside lower Physical Activity Level and Daily Steps. Insomnia is strongly associated with shorter Sleep Duration, lower Quality of Sleep, and elevated Stress Level. Age shows only a weak relationship with disorder status.

The same comparison was repeated for the three categorical features to test for demographic or occupational association with the target.




Figure 4.4b — Categorical Features vs. Sleep Disorder

Insights: BMI Category shows the clearest categorical association — Obese individuals skew heavily toward Sleep Apnea relative to Normal-weight individuals. Certain Occupations (Nurse, Doctor, Sales Representative) show elevated disorder rates, plausibly as a proxy for stress and shift-work patterns, while Gender shows minimal separation between classes.

4.5 Correlation Analysis

A Pearson correlation matrix of the numerical features was computed to surface linear relationships and flag potential multicollinearity ahead of model training.




Figure 4.5 — Correlation Matrix of Numerical Features

Insights: Sleep Duration and Quality of Sleep are strongly positively correlated ($r = 0.69$), as expected for two related sleep metrics. Physical Activity Level and Heart Rate are strongly negatively correlated ($r = -0.70$), and Systolic BP / Diastolic BP are strongly positively correlated ($r = 0.70$), confirming these are related but non-redundant cardiovascular signals. No pair reaches the near-perfect correlation that would force feature removal, so all numerical predictors were retained for modelling.

4.6 Multivariate Relationship Analysis

A pair-plot of the most target-relevant numerical features (Sleep Duration, Quality of Sleep, Stress Level, Heart Rate) was generated, coloured by class, to visually inspect how well these signals separate the three sleep-disorder categories in combination.



Figure 4.6 — Pair Plot of Key Features by Sleep Disorder Class

Insights: Insomnia (blue) and Sleep Apnea (green) each occupy visually distinct regions of the Sleep Duration $\times$ Heart Rate and Quality of Sleep $\times$ Heart Rate planes, while No Sleep Disorder (orange) dominates the moderate-stress, moderate-heart-rate region — evidence that even simple, linearly-separable structure exists in the feature space ahead of formal modelling.

4.7 Outlier Analysis

Outliers in the nine numerical features were assessed using the Interquartile Range (IQR) method and visualised with box plots.




Figure 4.7 — Box Plot Analysis of Numerical Features (Outlier Detection)

Key Insights: Diastolic BP shows the highest proportion of statistical outliers, followed by Sleep Duration and Age, each still representing under 1% of observations. These points correspond to physiologically plausible individuals (e.g. very low resting heart rate, unusually high blood pressure) rather than data-entry errors, so they were retained rather than removed — consistent with the objective of modelling real-world variation, including at-risk edge cases.

EDA Summary

Target is imbalanced (78% / 12% / 10%) → Macro F1 and class weighting/SMOTE required.

Sleep Apnea is driven by elevated blood pressure and heart rate; Insomnia by lower sleep quality, shorter duration, and higher stress.

BMI Category (Obese) is the strongest categorical signal; Gender carries little predictive value.

No feature pair is collinear enough to warrant removal → all predictors retained.

Outliers are limited to genuine physiological variation (<1% of rows) and were kept, not removed.

5. Feature Engineering

Three domain-informed features were derived from existing predictors to give the models additional, clinically meaningful signal without introducing any dependency on the target variable.

5.1 Engineered Features

Feature	Formula	Rationale
Pulse Pressure	Systolic BP – Diastolic BP	Additional cardiovascular signal beyond raw BP
Sleep Quality Index	$((\text{Sleep Duration} / 24) + (\text{Quality of Sleep} / 10)) / 2$	Consolidated, standardised sleep-profile score
Stress-Sleep Interaction	Stress Level $\times$ Sleep Duration	Captures the combined effect of stress and sleep duration

5.2 Engineered Feature Validation

Each engineered feature was checked for missing values (none were introduced) and validated by comparing its group-wise mean across the three target classes.

- Pulse Pressure shows strong separation for Sleep Apnea — mean 66.5 vs. approximately 49 for the other two classes — directly reflecting the elevated blood pressure observed during EDA.
- Sleep Quality Index is lowest for Insomnia (0.368) compared to approximately 0.50 for the other classes, consolidating two related sleep metrics into a single, standardised signal.
- Stress-Sleep Interaction captures individuals experiencing the combined burden of high stress and short sleep duration, a pattern that raw stress or raw sleep-duration alone does not fully express.

All three engineered features separate the target classes meaningfully and were retained for modelling.

6. Data Preprocessing & Train-Test Split

6.1 Preprocessing Pipeline

All preprocessing was implemented as a single scikit-learn ColumnTransformer, fit exclusively on the training fold in every cross-validation split to prevent information leakage from the test set.

Feature Type	Pipeline
Numerical	Median imputation → StandardScaler
Categorical	Mode imputation → OneHotEncoder

Combining the numerical and one-hot-encoded categorical outputs expands the processed feature matrix from 13 raw columns to 30 model-ready features.

6.2 Stratified Train-Test Split

Given the 78% / 12% / 10% class imbalance, a random split risks under-representing Insomnia or Sleep Apnea in either partition. A stratified 80/20 split was therefore used to preserve the original class proportions in both the training and test sets.

Split	Samples	Class Proportions Preserved
Training Set	8,000	Yes (stratified)
Test Set	2,000	Yes (stratified)

The test set remained completely untouched until the final evaluation stage in Section 8.

7. Model Development & Selection

7.1 Modeling Approach

Five classification algorithms — Logistic Regression, Decision Tree, Random Forest, Support Vector Machine, and XGBoost — were each benchmarked under two imbalance-handling strategies, class weighting and SMOTE (Synthetic Minority Over-sampling), using Stratified 5-Fold Cross-Validation. Preprocessing and SMOTE (where applicable) were refit independently inside every fold to guarantee a leakage-safe comparison.

Primary metric: Macro F1, which weights every class equally regardless of size — the appropriate choice for a 78/12/10 imbalanced target. Balanced Accuracy, Macro Precision/Recall, and ROC-AUC were tracked as supporting evidence.

7.2 Baseline Cross-Validation Results

Class-Weighting Strategy

Model	Accuracy	Balanced Acc.	F1 Macro	ROC-AUC
XGBoost	91.58%	81.88%	84.38%	96.79%
Random Forest	90.58%	76.66%	81.63%	—
SVM	87.35%	88.17%	80.71%	—
Logistic Regression	86.69%	88.73%	80.11%	—
Decision Tree	86.05%	74.56%	75.02%	—

SMOTE Strategy

Model	Accuracy	Balanced Acc.	F1 Macro
XGBoost + SMOTE	91.35%	82.92%	84.28%
Random Forest + SMOTE	90.74%	82.62%	83.36%
SVM + SMOTE	88.14%	86.95%	81.35%
Logistic Regression + SMOTE	87.51%	88.65%	81.00%
Decision Tree + SMOTE	85.13%	76.29%	74.54%

XGBoost led under both strategies with near-identical Macro F1 (84.38% vs. 84.28%), and was carried forward — alongside Random Forest, its closest competitor — into hyperparameter tuning.

7.3 Class Weight vs. SMOTE Comparison

The two imbalance-handling strategies were compared directly across all five candidate models on cross-validated Macro F1.




Figure 7.3 — Cross-Validated Macro F1: Class Weight vs. SMOTE

Observation: the two strategies perform comparably across every model, with SMOTE holding a slight edge for Random Forest, SVM, and Logistic Regression, while class weighting is marginally stronger for XGBoost and Decision Tree. Given the near-parity in performance, class weighting was preferred going forward — it avoids generating synthetic samples and is simpler to reproduce consistently at inference time.

7.4 Hyperparameter Tuning

The strongest candidates from the baseline comparison — XGBoost (both strategies) and Random Forest — were tuned using RandomizedSearchCV with Stratified 5-Fold cross-validation, optimising for Macro F1. The test set remained untouched throughout this stage.

Candidate	Best CV Macro F1
XGBoost — Class Weight (selected)	85.33%
XGBoost — SMOTE	85.05%
Random Forest — Tuned	83.26%

Selected Configuration — XGBoost (Class Weight)

Hyperparameter	Tuned Value
n_estimators	500
max_depth	3
learning_rate	0.05
subsample	1.0
colsample_bytree	1.0
min_child_weight	1

SELECTED — XGBoost (Class Weight, tuned) selected as the final model — highest cross-validated Macro F1 among all candidates tested.

8. Final Model Evaluation

The tuned XGBoost model was evaluated once, on the 2,000-record held-out test set that had not been used during training, tuning, or cross-validation — giving a realistic estimate of production performance.

8.1 Held-Out Test Set Performance

Metric	Score
Accuracy	93.50%
Balanced Accuracy	87.43%
F1 Macro	88.50%
F1 Weighted	93.42%
Precision Macro	89.65%
Recall Macro	87.43%
ROC-AUC (One-vs-Rest)	98.29%
PR-AUC (Macro Avg., One-vs-Rest)	94.88%

The test Macro F1 (88.50%) exceeds the cross-validated average (85.33%), and overall accuracy (93.50%) comfortably clears the 78% no-skill baseline of always predicting the majority class — strong evidence of genuine generalisation

rather than overfitting to the training folds.

8.2 Confusion Matrix

The confusion matrix breaks down prediction outcomes by true class to reveal exactly where the model succeeds and where it confuses classes.




Figure 8.2 — Final XGBoost Confusion Matrix (Test Set)

Key Insights: the model correctly classifies 1,511 of 1,563 No Sleep Disorder cases and 167 of 190 Sleep Apnea cases. Insomnia is the hardest class to isolate — 54 of 247 true Insomnia cases are misclassified as No Sleep Disorder, while confusion between Insomnia and Sleep Apnea is minimal (at most 3 records each direction), showing the model rarely confuses the two disorder types with each other — only with the healthy baseline.

8.3 Multiclass ROC-AUC

One-vs-Rest ROC curves were plotted for each class to evaluate ranking quality independent of the classification threshold.




Figure 8.3 — Multiclass ROC-AUC Curve, Final XGBoost Model

Insights: all three classes achieve AUC above 0.97 — No Sleep Disorder 0.976, Insomnia 0.980, and Sleep Apnea 0.993 — indicating excellent discrimination across every decision threshold, not just the default 50% cut-off used for the confusion matrix above.

8.4 Precision-Recall Curves & PR-AUC

Because the target is imbalanced, Precision-Recall analysis was used alongside ROC-AUC to give a second, threshold-independent view of class-specific performance. One-vs-Rest PR curves were plotted for each class using the final XGBoost model, with particular attention to the more difficult, less-separated classes.




Figure 8.4 — Multiclass Precision-Recall Curves, Final XGBoost Model

Insights: PR-AUC is 0.993 for No Sleep Disorder and 0.959 for Sleep Apnea — both comfortably above 0.95 — while Insomnia trails at 0.895. Because PR-AUC is more sensitive than ROC-AUC to performance on the minority classes, this gap reinforces the confusion-matrix finding that Insomnia is the model's hardest class to separate from a healthy profile, even though its ROC-AUC alone (0.980) looked strong.

8.5 Class-wise Performance

Precision, Recall, and F1-score were computed independently for each class to identify which disorder is hardest for the model to detect.

Class	Precision	Recall	F1-Score
No Sleep Disorder	95.33%	96.67%	96.00%
Sleep Apnea	89.78%	87.89%	88.83%
Insomnia	83.84%	77.73%	80.67%




Figure 8.5 — Class-wise Precision, Recall, and F1-Score

Reading the results: No Sleep Disorder and Sleep Apnea are both detected reliably (F1 above 88%), while Insomnia remains the hardest class, most often confused with a healthy profile — consistent with the EDA finding that Insomnia's physiological signals overlap more with No Sleep Disorder than Sleep Apnea's more distinctive cardiovascular profile does.

9. Model Explainability (XAI)

To ensure the final XGBoost model is not a 'black box', SHAP (SHapley Additive exPlanations) was applied via TreeExplainer on the transformed test set to quantify how much each feature contributes to every prediction, both globally and by feature interaction.

9.1 SHAP Global Feature Importance

Aggregating SHAP values across all test observations, split by target class, produces a global importance ranking for every feature in the model.




Figure 9.1 — SHAP Global Feature Importance by Target Class

Insights: Heart Rate is the single strongest predictor across all classes, followed by Sleep Duration and the engineered Sleep Quality Index. Systolic BP and Age contribute substantially, while Physical Activity Level and Pulse Pressure provide meaningful supporting signal. Demographic and occupational one-hot features (e.g. Occupation_Doctor, Occupation_Teacher) rank lowest overall — the model relies primarily on physiological and sleep-related signals rather than demographic proxies.

9.2 SHAP Feature Interaction Analysis

An interaction summary was generated for the model's top continuous predictors to examine whether pairs of features jointly influence predictions beyond their individual, additive effects.




Figure 9.2 — SHAP Feature Interaction Summary (Quality of Sleep, Sleep Duration, Age)

Key Insight: Sleep Duration shows the strongest interaction effect with Quality of Sleep — the impact of one depends visibly on the level of the other, reflecting their real-world relationship as two facets of the same underlying sleep-health construct rather than fully independent signals.

9.3 Explainability Summary

Key Insight

Heart Rate, Sleep Duration, and the Sleep Quality Index dominate the model's predictions across both the global SHAP ranking and the interaction analysis. Cardiovascular indicators (Systolic BP, Pulse Pressure) provide strong secondary signal, particularly for Sleep Apnea, while demographic and occupational features contribute comparatively little. This consistent, human-readable explanation is exactly what makes the model suitable for a real-world screening context, where 'why' matters as much as 'what'.

10. Conclusion & Key Takeaways

10.1 Summary of Findings

- XGBoost achieved the strongest and most consistent performance among all five algorithms tested, under both class-weighting and SMOTE imbalance strategies.
- Class weighting and SMOTE performed comparably across every model — class weighting was marginally better for the final model and simpler to deploy, since it requires no synthetic-sample generation at inference time.
- Heart Rate, Sleep Duration, and Sleep Quality Index are the strongest predictors of sleep-disorder status, confirmed independently by both coefficient-level and SHAP-based explainability.

- Sleep Apnea is driven primarily by elevated blood pressure and heart rate; Insomnia by lower sleep quality and duration combined with higher stress.
- Insomnia remains the hardest class to separate from a healthy sleep profile — the clearest target for future feature or data improvements.

10.2 Final Recommendation

The tuned XGBoost (Class Weight) model is recommended as the deployment candidate for this classification task, combining the best held-out test performance (93.50% accuracy, 88.50% Macro F1, 98.29% ROC-AUC) with full interpretability via SHAP — a combination essential for any model intended to support health-related screening decisions.

10.3 Business Impact & Recommendations

- Position the model as a screening / triage aid, not a diagnostic system — flagged cases should be routed to a qualified healthcare provider for confirmation.
- Prioritise Insomnia recall in future iterations, given it is the class most often confused with a healthy profile and therefore the highest-risk false negative.
- Use the SHAP-based explanations to give end users (clinicians or individuals) a transparent, feature-level rationale alongside every prediction, not just a class label.
- Validate the pipeline on real clinical data before any production use, since the current dataset is synthetically generated.

10.4 Skills & Techniques Demonstrated

- Data cleaning, validation, and exploratory data analysis (pandas, seaborn, matplotlib).
- Imbalanced-classification methodology: stratified splitting, class weighting, and SMOTE, compared empirically rather than assumed.
- Feature engineering for health and lifestyle data: pulse pressure, a composite sleep-quality index, and a stress-sleep interaction term.
- Leakage-safe pipeline design using scikit-learn's ColumnTransformer and Pipeline, with resampling refit inside every cross-validation fold.
- Model development across linear (Logistic Regression), tree-based (Decision Tree, Random Forest), kernel-based (SVM), and gradient-boosted (XGBoost) paradigms.
- Hyperparameter optimisation using RandomizedSearchCV with Stratified K-Fold cross-validation.
- Explainable AI (XAI) using SHAP TreeExplainer — global feature importance and feature-interaction analysis.
- Clear, structured technical communication — translating model internals into business-relevant conclusions for non-technical stakeholders.

— End of Report —